

## Assignment -1 Test Disk Performance IOPS With Fio utility for random read write operations

### Instructions

Run this command:

```
fio --randrepeat=1 --ioengine=libaio --direct=1 --gtod_reduce=1 --name=test
--filename=test
--bs=2k --iodepth=64 --size=2G --readwrite=randrw --rwmixread=75
```

A:

```
[pranay@localhost ~]$ fio --randrepeat=1 --ioengine=libaio --direct=1 --gtod_reduce=1 --name=test --filename=test --bs=2
4 --size=2G --readwrite=randrw --rwmixread=75
test: (g=0): rw=randrw, bs=(R) 2048B-2048B, (W) 2048B-2048B, (T) 2048B-2048B, ioengine=libaio, iodepth=64
fio-3.7
Starting 1 process
test: Laying out IO file (1 file / 2048MiB)
Jobs: 1 (f=1): [m(1)][99.6%][r=7249KiB/s,w=2532KiB/s][r=3624,w=1266 IOPS][eta 00m:01s]
test: (groupid=0, jobs=1): err= 0: pid=1500: Wed Aug 23 14:31:31 2023
  read: IOPS=3048, BW=6097KiB/s (6243kB/s)(1535MiB/257815msec)
    bw ( KiB/s): min= 609, max= 8100, per=99.91%, avg=6090.22, stdev=1225.45, samples=515
    iops        : min= 304, max= 4050, avg=3045.05, stdev=612.72, samples=515
  write: IOPS=1018, BW=2038KiB/s (2086kB/s)(513MiB/257815msec)
    bw ( KiB/s): min= 209, max= 2772, per=99.92%, avg=2035.28, stdev=421.58, samples=515
    iops        : min= 104, max= 1386, avg=1017.57, stdev=210.79, samples=515
  cpu      : usr=1.40%, sys=33.51%, ctx=412355, majf=2, minf=24
IO depths  : 1=0.1%, 2=0.1%, 4=0.1%, 8=0.1%, 16=0.1%, 32=0.1%, >=64=100.0%
 submit    : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%, 64=0.0%, >=64=0.0%
 complete  : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%, 64=0.1%, >=64=0.0%
 issued rwts: total=785920,262656,0,0 short=0,0,0,0 dropped=0,0,0,0
 latency   : target=0, window=0, percentile=100.00%, depth=64

Run status group 0 (all jobs):
  READ: bw=6097KiB/s (6243kB/s), 6097KiB/s-6097KiB/s (6243kB/s-6243kB/s), io=1535MiB (1610MB), run=257815-257815msec
  WRITE: bw=2038KiB/s (2086kB/s), 2038KiB/s-2038KiB/s (2086kB/s-2086kB/s), io=513MiB (538MB), run=257815-257815msec

Disk stats (read/write):
  sda: ios=800939/262412, merge=5362/6, ticks=498932/60665, in_queue=558841, util=92.95%
```

—readwrite=randwrite

```
[pranay@localhost ~]$ fio --randrepeat=1 --ioengine=libaio --direct=1 --gtod_reduce=1 --name=test --filename=test --bs=2k
4 --size=2G --readwrite=randwrite --rwmixread=75
test: (g=0): rw=randwrite, bs=(R) 2048B-2048B, (W) 2048B-2048B, (T) 2048B-2048B, ioengine=libaio, iodepth=64
fio-3.7
Starting 1 process
Jobs: 1 (f=1): [w(1)][100.0%][r=0KiB/s,w=7771KiB/s][r=0,w=3885 IOPS][eta 00m:00s]
test: (groupid=0, jobs=1): err= 0: pid=1748: Wed Aug 23 14:57:47 2023
  write: IOPS=3696, BW=7394KiB/s (7571kB/s)(2048MiB/283638msec)
    bw ( KiB/s): min= 590, max=10948, per=99.97%, avg=7391.04, stdev=1675.47, samples=567
    iops        : min= 295, max= 5474, avg=3695.49, stdev=837.74, samples=567
  cpu      : usr=2.12%, sys=6.50%, ctx=1048583, majf=1, minf=27
IO depths  : 1=0.1%, 2=0.1%, 4=0.1%, 8=0.1%, 16=0.1%, 32=0.1%, >=64=100.0%
 submit    : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%, 64=0.0%, >=64=0.0%
 complete  : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%, 64=0.1%, >=64=0.0%
 issued rwts: total=0,1048576,0,0 short=0,0,0,0 dropped=0,0,0,0
 latency   : target=0, window=0, percentile=100.00%, depth=64

Run status group 0 (all jobs):
  WRITE: bw=7394KiB/s (7571kB/s), 7394KiB/s-7394KiB/s (7571kB/s-7571kB/s), io=2048MiB (2147MB), run=283638-283638msec

Disk stats (read/write):
  sda: ios=49940/1049288, merge=7485/23712, ticks=236526/243706, in_queue=479975, util=85.04%
```

—readwrite=randread

```
[pranay@localhost ~]$ fio --randrepeat=1 --ioengine=libaio --direct=1 --gtod_reduce=1 --name=test --filename=test --bs=2k
4 --size=2G --readwrite=randread --rwmixread=75
test: (groupid=0, jobs=1): err= 0: pid=1822: Wed Aug 23 15:03:04 2023
fio-3.7
Starting 1 process
Jobs: 1 (f=1): [r(1)][100.0%][r=9.78MiB/s,w=0KiB/s][r=5008,w=0 IOPS][eta 00m:00s]
test: (groupid=0, jobs=1): err= 0: pid=1822: Wed Aug 23 15:03:04 2023
read: IOPS=4185, BW=8370KiB/s (8571kB/s)(2048MiB/250554msec)
bw ( KiB/s): min= 277, max= 5278, per=21.54%, avg=1803.30, stdev=469.56, samples=496
iops: min= 138, max= 2639, avg=901.40, stdev=234.79, samples=496
cpu: usr=4.06%, sys=58.96%, ctx=367760, majf=2, minf=96
IO depths : 1=0.1%, 2=0.1%, 4=0.1%, 8=0.1%, 16=0.1%, 32=0.1%, >=64=100.0%
submit : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%, 64=0.0%, >=64=0.0%
complete : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%, 64=0.1%, >=64=0.0%
issued rwts: total=1048576,0,0,0 short=0,0,0,0 dropped=0,0,0,0
latency : target=0, window=0, percentile=100.00%, depth=64

Run status group 0 (all jobs):
READ: bw=8370KiB/s (8571kB/s), 8370KiB/s-8370KiB/s (8571kB/s-8571kB/s), io=2048MiB (2147MB), run=250554-250554msec

Disk stats (read/write):
sda: ios=1066637/1770, merge=14188/49211, ticks=15125128/31139, in_queue=15153926, util=99.97%
You have new mail in /var/spool/mail/pranay
```

## Assignment 2- Setup NFS server in one linux system and access the shared directory from another linux system which will act as NFS client

Refer Network file system slide for this Assignment

### A: SERRVER SIDE packages installation

```
[pranay@localhost ~]$ sudo yum install nfs-utils libnfsidmap -y
[sudo] password for pranay:
Loaded plugins: fastestmirror
Loading mirror speeds from cached hostfile
* base: centos.excellmedia.net
* extras: centos.excellmedia.net
* updates: centos.excellmedia.net
base | 3.6 kB | 00:00:00
extras | 2.9 kB | 00:00:00
mysql-connectors-community | 2.6 kB | 00:00:00
mysql-tools-community | 2.6 kB | 00:00:00
mysql57-community | 2.6 kB | 00:00:00
nginx | 2.9 kB | 00:00:00
updates | 2.9 kB | 00:00:00
updates/7/x86_64/primary_db | 22 MB | 00:00:04
Package libnfsidmap-0.25-19.el7.x86_64 already installed and latest version
Resolving Dependencies
--> Running transaction check
---> Package nfs-utils.x86_64 1:1.3.0-0.66.el7 will be updated
---> Package nfs-utils.x86_64 1:1.3.0-0.68.el7.2 will be an update
--> Finished Dependency Resolution
```

### Services status

```
[pranay@localhost ~]$ sudo systemctl status rpcbind.service nfs-server.service rpc-statd.service nfs-idmapd.service
● rpcbind.service - RPC bind service
   Loaded: loaded (/usr/lib/systemd/system/rpcbind.service; enabled; vendor preset: enabled)
   Active: active (running) since Wed 2023-08-30 10:56:11 IST; 3h 46min ago
   Main PID: 365 (rpcbind)
   CGroup: /system.slice/rpcbind.service
           └─365 /sbin/rpcbind -w

Aug 30 10:56:10 localhost.localdomain systemd[1]: Starting RPC bind service...
Aug 30 10:56:11 localhost.localdomain systemd[1]: Started RPC bind service.

● nfs-server.service - NFS server and services
   Loaded: loaded (/usr/lib/systemd/system/nfs-server.service; enabled; vendor preset: disabled)
   Active: active (exited) since Wed 2023-08-30 14:42:41 IST; 27s ago
   Process: 3424 ExecStartPost=/bin/sh -c if systemctl -q is-active gssproxy; then systemctl reload gssproxy ; fi (code=exited, status=0/SUCCESS)
   Process: 3408 ExecStart=/usr/sbin/rpc.nfsd $RPCNFSDARGS (code=exited, status=0/SUCCESS)
   Process: 3407 ExecStartPre=/usr/sbin/exportfs -r (code=exited, status=0/SUCCESS)
   Main PID: 3408 (code=exited, status=0/SUCCESS)
   CGroup: /system.slice/nfs-server.service

Aug 30 14:42:41 localhost.localdomain systemd[1]: Starting NFS server and services...
Aug 30 14:42:41 localhost.localdomain systemd[1]: Started NFS server and services.

● rpc-statd.service - NFS status monitor for NFSv2/3 locking.
   Loaded: loaded (/usr/lib/systemd/system/rpc-statd.service; static; vendor preset: disabled)
   Active: active (running) since Wed 2023-08-30 14:42:41 IST; 28s ago
   Process: 3397 ExecStart=/usr/sbin/rpc.statd $STATDARGS (code=exited, status=0/SUCCESS)
```

## Directories and packages

```
[pranay@localhost ~]$ sudo mkdir server
[sudo] password for pranay:
Sorry, try again.
[sudo] password for pranay:
[pranay@localhost ~]$ cd server
[pranay@localhost server]$ sudo mkdir apps
[pranay@localhost server]$ ls
apps
[pranay@localhost server]$ chmod 777 server
chmod: cannot access 'server': No such file or directory
[pranay@localhost server]$ cd ..
You have new mail in /var/spool/mail/pranay
[pranay@localhost ~]$ chmod 777 server
chmod: changing permissions of 'server': Operation not permitted
[pranay@localhost ~]$ sudo chmod 777 /server
[pranay@localhost ~]$ sudo chmod 777 /apps
chmod: cannot access '/apps': No such file or directory
[pranay@localhost ~]$ sudo chmod 777 /server/apps
```

## Nfs exports

```
[pranay@localhost server]$ cat /etc/exports
/server/apps 192.168.56.103(rw,sync,no_root_squash)
[pranay@localhost server]$ exports -rv
-bash: exports: command not found
[pranay@localhost server]$ exportfs -rv
exporting 192.168.56.103:/server/apps
exportfs: could not open /var/lib/nfs/.etab.lock for locking: errno 13 (Permission denied)
exportfs: can't lock /var/lib/nfs/etab for writing
[pranay@localhost server]$ sudo exportfs -rv
[sudo] password for pranay:
exporting 192.168.56.103:/server/apps
```

## CLIENT SIDE

### Packages installation

```
[pranay@localhost ~]$ sudo yum install nfs-utils rpcbind -y
[sudo] password for pranay:
Loaded plugins: fastestmirror
Determining fastest mirrors
 * base: mirrors.nhanhoa.com
 * extras: mirror.xtom.com.hk
 * updates: mirror.xtom.com.hk
base | 3.6 kB 00:00:00
extras | 2.9 kB 00:00:00
mysql-connectors-community | 2.6 kB 00:00:00
mysql-tools-community | 2.6 kB 00:00:00
mysql57-community | 2.6 kB 00:00:00
nginx | 2.9 kB 00:00:00
updates | 2.9 kB 00:00:00
(1/6): mysql57-community/x86_64/primary_db | 349 kB 00:00:00
(2/6): extras/7/x86_64/primary_db | 250 kB 00:00:00
(3/6): mysql-connectors-community/x86_64/primary_db | 100 kB 00:00:00
(4/6): mysql-tools-community/x86_64/primary_db | 94 kB 00:00:00
(5/6): nginx/x86_64/primary_db | 85 kB 00:00:02
```

### Client rpc services

```
[pranay@localhost ~]$ sudo systemctl enable rpcbind
[pranay@localhost ~]$ sudo systemctl start rpcbind
[pranay@localhost ~]$ sudo systemctl status rpcbind
● rpcbind.service - RPC bind service
   Loaded: loaded (/usr/lib/systemd/system/rpcbind.service; enabled; vendor preset: enabled)
   Active: active (running) since Wed 2023-08-30 09:45:46 UTC; 3min 32s ago
   Main PID: 399 (rpcbind)
   CGroup: /system.slice/rpcbind.service
           └─399 /sbin/rpcbind -w

Aug 30 09:45:46 localhost.localdomain systemd[1]: Starting RPC bind service...
Aug 30 09:45:46 localhost.localdomain systemd[1]: Started RPC bind service.
```

## Mounting client

```
[pranay@localhost ~]$ sudo systemctl stop firewalld.service
[pranay@localhost ~]$ showmount -e 192.168.56.101
Export list for 192.168.56.101:
/server/apps 192.168.56.103

[pranay@localhost ~]$ sudo mkdir -p /mnt/apps
[pranay@localhost ~]$ cd mnt
-bash: cd: mnt: No such file or directory
[pranay@localhost ~]$ cd /mnt
[pranay@localhost mnt]$ ls
apps
[pranay@localhost mnt]$ cd ..
[pranay@localhost /]$ mount 192.168.56.101:/server/apps /mnt/apps
mount: only root can do that
[pranay@localhost /]$ sudo mount 192.168.56.101:/server/apps /mnt/apps
```

## Output

### serverside

```
[pranay@localhost server]$ cd apps
[pranay@localhost apps]$ ls
file
[pranay@localhost apps]$ cat file
i have done nfs assignment and the output now u r seeing is from the server side
but this is created on the client side
```

### Client side

```
[pranay@localhost /]$ cd mnt
[pranay@localhost mnt]$ cd apps
[pranay@localhost apps]$ ls
[pranay@localhost apps]$ touch file
[pranay@localhost apps]$ ls
file
[pranay@localhost apps]$ vi file
```


**Assignment 3 - Add new Virtual hard disk Drive for the virtual box and partition the drive in linux then create file system and mount the file system**

**Note:**

Take a clone of your machine before starting this assignment

### Instructions

1. select the Virtual Box, Click on Settings -> Storage, select hard drive and click on Add a hard disk.
2. Create a new disk following the on-screen instructions. Provide the directory on windows where you want the drive created.



## Partition the new drive

### 1. Run "lsblk" command to check the partitions

A:

```
[pranay@localhost ~]$ lsblk
NAME        MAJ:MIN RM  SIZE RO TYPE MOUNTPOINT
sda          8:0    0   40G  0 disk
└─sda1       8:1    0   40G  0 part /
sdb          8:16   0   40G  0 disk
```

### 2. create a new partition using the command "fdisk /dev/sdb"

A:

```
[pranay@localhost ~]$ sudo fdisk -l

Disk /dev/sda: 42.9 GB, 42949672960 bytes, 83886080 sectors
Units = sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
Disk label type: dos
Disk identifier: 0x0009ef1a

   Device Boot      Start         End      Blocks   Id  System
/dev/sda1    *        2048     83886079     41942016   83   Linux

Disk /dev/sdb: 42.9 GB, 42949672960 bytes, 83886080 sectors
Units = sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
```

### 3. p option will show the list of partitions

### 4. n option will create new partition

A:

```
Command (m for help): n
Partition type:
   p   primary (0 primary, 0 extended, 4 free)
   e   extended
Select (default p):
Using default response p
Partition number (1-4, default 1):
First sector (2048-83886079, default 2048):
Using default value 2048
Last sector, +sectors or +size{K,M,G} (2048-83886079, default 83886079):
Using default value 83886079
Partition 1 of type Linux and of size 40 GiB is set

Command (m for help): w
The partition table has been altered!

Calling ioctl() to re-read partition table.
Syncing disks.
```

### 5. Create 1 primary partition

### 6. specify where the partition will begin and end by accepting the default values

### 7. w option will save the changes



(Note: Since this is the first partition we can start at the first available sector and since we want to use the entire disk we can specify the last sector as the end. Note that, if you wish to create multiple partitions - you can even specify the size of each partition

by sectors, bytes, kilobytes or megabytes)

8. Create the file system on the partition using the command `"/usr/sbin/mkfs.ext4 -L data /dev/sdb1"`

A:

```
[pranay@localhost ~]$ sudo /usr/sbin/mkfs.ext4 -L data /dev/sdb1
mke2fs 1.42.9 (28-Dec-2013)
Filesystem label=data
OS type: Linux
Block size=4096 (log=2)
Fragment size=4096 (log=2)
Stride=0 blocks, Stripe width=0 blocks
2621440 inodes, 10485504 blocks
524275 blocks (5.00%) reserved for the super user
First data block=0
Maximum filesystem blocks=2157969408
320 block groups
32768 blocks per group, 32768 fragments per group
8192 inodes per group
Superblock backups stored on blocks:
    32768, 98304, 163840, 229376, 294912, 819200, 884736, 1605632, 2654208,
    4096000, 7962624

Allocating group tables: done
Writing inode tables: done
Creating journal (32768 blocks): done
Writing superblocks and filesystem accounting information: done
```

9. Create the directory `/data`

10. Mount the file system using the command `"mount LABEL=data /data"`

A:

```
[pranay@localhost ~]$ sudo mkdir /data
[pranay@localhost ~]$ sudo mount LABEL=data /data
```

Run the mount command with no arguments which shows all currently mounted filesystems. You can use the below command as well

`lsblk -o NAME,FSTYPE,SIZE,MOUNTPOINT,LABEL`

```
[pranay@localhost ~]$ lsblk -o NAME,FSTYPE,SIZE,MOUNTPOINT,LABEL
NAME    FSTYPE  SIZE MOUNTPOINT LABEL
sda                        40G
└─sda1 xfs      40G /
sdb                        40G
└─sdb1 ext4    40G /data    data
```

11. To automatically mount the file system at the time of boot we need the below entry to be added to the `/etc/fstab` file

`LABEL=data /data ext4 defaults 12`

A:

```
[pranay@localhost ~]$ sudo cat /etc/fstab
#
# /etc/fstab
# Created by anaconda on Thu Apr 30 22:04:55 2020
#
# Accessible filesystems, by reference, are maintained under '/dev/disk'
# See man pages fstab(5), findfs(8), mount(8) and/or blkid(8) for more info
#
UUID=1c419d6c-5064-4a2b-953c-05b2c67edb15 / xfs defaults 0 0
/swapfile none swap defaults 0 0
LABEL=data /data ext4 defaults 12
```

## Assignment 4 - Expanding a filesystem on the fly using Logical Volume Management

### Instructions

Note: install LVM using `sudo yum install lvm2`

### Adding a new logical volume

#### 1.1 Install a new hard drive.

A:

```
[pranay@localhost ~]$ sudo yum install lvm2
[sudo] password for pranay:
Loaded plugins: fastestmirror
Loading mirror speeds from cached hostfile
 * base: centos.excellmedia.net
 * extras: centos.excellmedia.net
 * updates: centos.excellmedia.net
base | 3.6 kB 00:00:00
extras | 2.9 kB 00:00:00
mysql-connectors-community | 2.6 kB 00:00:00
mysql-tools-community | 2.6 kB 00:00:00
mysql57-community | 2.6 kB 00:00:00
nginx | 2.9 kB 00:00:00
updates | 2.9 kB 00:00:00
Resolving Dependencies
--> Running transaction check
--> Package lvm2.x86_64 7:2.02.187-6.el7_9.5 will be installed
--> Processing Dependency: lvm2-libs = 7:2.02.187-6.el7_9.5 for package: 7:lvm2-2.02.187-6.el7_9.5.x86_64
--> Processing Dependency: device-mapper-persistent-data >= 0.7.0-0.1.rc6 for package: 7:lvm2-2.02.187-6.el7_9.5.x86_64
--> Processing Dependency: liblvm2app.so.2.2(Base)(64bit) for package: 7:lvm2-2.02.187-6.el7_9.5.x86_64
--> Processing Dependency: libdevmapper-event.so.1.02(Base)(64bit) for package: 7:lvm2-2.02.187-6.el7_9.5.x86_64
--> Processing Dependency: liblvm2app.so.2.2()(64bit) for package: 7:lvm2-2.02.187-6.el7_9.5.x86_64
--> Processing Dependency: libdevmapper-event.so.1.02()(64bit) for package: 7:lvm2-2.02.187-6.el7_9.5.x86_64
--> Running transaction check
--> Package device-mapper-event-libs.x86_64 7:1.02.170-6.el7_9.5 will be installed
--> Package device-mapper-persistent-data.x86_64 0:0.8.5-3.el7_9.2 will be installed
--> Package lvm2-libs.x86_64 7:2.02.187-6.el7_9.5 will be installed
--> Processing Dependency: device-mapper-event = 7:1.02.170-6.el7_9.5 for package: 7:lvm2-libs-2.02.187-6.el7_9.5.x86_64--> Running transaction check
```

#### 1.2 Use `lsblk` command to identify the new hard drive.

#### 1.3 Create Physical Volume from the hard drive using "`pvcreate /dev/sdc`". Assuming `/dev/sdc` is the new hard drive created.

#### 1.4 You can list all the PV with the command:`pvscan`

#### 1.5 If you want to display more information about any specific PV, let's say `/dev/sdc` run

`"pvdisplay/dev/sdc"`

#### 1.6 Create Volume Group named "share" using `sudo vgcreate share /dev/sdc`

## A: 1.2 to 1.6

```
pranay@localhost:~  
Last login: Sat Sep 2 10:56:40 2023  
[pranay@localhost ~]$ lsblk  
NAME MAJ:MIN RM SIZE RO TYPE MOUNTPOINT  
sda 8:0 0 40G 0 disk  
└─sda1 8:1 0 40G 0 part /  
sdb 8:16 0 40G 0 disk  
└─sdb1 8:17 0 40G 0 part /data  
sdc 8:32 0 40G 0 disk  
[pranay@localhost ~]$ sudo pvcreate /dev/sdc  
[sudo] password for pranay:  
Physical volume "/dev/sdc" successfully created.  
[pranay@localhost ~]$ pvscan  
WARNING: Running as a non-root user. Functionality may be unavailable.  
/run/lvm/lvmetad.socket: access failed: Permission denied  
WARNING: Failed to connect to lvmetad. Falling back to device scanning.  
/run/lock/lvm/P_global:aux: open failed: Permission denied  
Unable to obtain global lock.  
[pranay@localhost ~]$ sudo pvscan  
PV /dev/sdc lvm2 [40.00 GiB]  
Total: 1 [40.00 GiB] / in use: 0 [0 ] / in no VG: 1 [40.00 GiB]  
[pranay@localhost ~]$ sudo pvdisplay /dev/sdc  
"/dev/sdc" is a new physical volume of "40.00 GiB"  
--- NEW Physical volume ---  
PV Name /dev/sdc  
VG Name  
PV Size 40.00 GiB  
Allocatable NO  
PE Size 0  
Total PE 0  
Free PE 0  
Allocated PE 0  
PV UUID 0oPM2H-4680-sH7o-jXn8-2T3W-XUga-iWAn99  
[pranay@localhost ~]$ sudo vgcreate share /dev/sdc  
Volume group "share" successfully created  
You have new mail in /var/spool/mail/pranay
```

1.7 you can list all the VGs with the command:vgscan

1.8 You can display more information about any specific VG, such as share with the command:vgdisplay share

1.9 Create Logical Volume: To create a 100MB LV mylv from VG share use the command

"sudo lvcreate --size 100M --name mylv share"

1.10 create another LV mylv1 of size 1GB from VG share with the command:"sudo lvcreate

--size 1G --name mylv1 share"

1.11 list all the LVs with the commands: lvscan or lvs



## A: 1.7 to 1.11

```
[pranay@localhost ~]$ sudo vgscan
  Reading volume groups from cache.
  Found volume group "share" using metadata type lvm2
[pranay@localhost ~]$ sudo vgdisk share
--- Volume group ---
VG Name          share
System ID
Format           lvm2
Metadata Areas   1
Metadata Sequence No 1
VG Access        read/write
VG Status        resizable
MAX LV          0
Cur LV         0
Open LV         0
Max PV          0
Cur PV         1
Act PV         1
VG Size         <40.00 GiB
PE Size         4.00 MiB
Total PE        10239
Alloc PE / Size 0 / 0
Free PE / Size  10239 / <40.00 GiB
VG UUID         QT0d6y-17EA-89Z0-E2Ns-jemF-nPpc-bGfYBZ

[pranay@localhost ~]$ sudo lvcreate --size 100M --name mylv share
  Logical volume "mylv" created.
[pranay@localhost ~]$ sudo lvcreate --size 1G --name mylv1 share
  Logical volume "mylv1" created.
[pranay@localhost ~]$ sudo lvs
LV      VG      Attr      LSize   Pool Origin Data%  Meta%  Move Log Cpy%Sync Convert
mylv    share -wi-a----- 100.00m
mylv1   share -wi-a-----  1.00g

[pranay@localhost ~]$ sudo lvdisplay mylv
```

**1.12** You can also display more information about any specific LV with the command: "sudo lvdisplay VG\_NAME/LV\_NAME"

**1.13** The LVs are available as /dev/VG\_NAME/LV\_NAME. You should be able to see the 2 LVs created so far.

## A: 1.12 to 1.13

```
pranay@localhost:~$ sudo lvsdisplay share
--- Logical volume ---
LV Path                /dev/share/mylv
LV Name                 mylv
VG Name                 share
LV UUID                S02dNj-AbLH-FAKm-Ve9l-rGRp-7u3E-pfzQ9m
LV Write Access         read/write
LV Creation host, time  localhost.localdomain, 2023-09-02 11:06:54 +0530
LV Status                available
# open                  0
LV Size                 100.00 MiB
Current LE               25
Segments                1
Allocation              inherit
Read ahead sectors      auto
- currently set to      8192
Block device            253:0

--- Logical volume ---
LV Path                /dev/share/mylv1
LV Name                 mylv1
VG Name                 share
LV UUID                qx6tcJ-bx2Q-0KWv-F2G0-cOYD-SwKf-kd5T4q
LV Write Access         read/write
LV Creation host, time  localhost.localdomain, 2023-09-02 11:08:09 +0530
LV Status                available
# open                  0
LV Size                 1.00 GiB
Current LE               256
Segments                1
Allocation              inherit
Read ahead sectors      auto
- currently set to      8192
Block device            253:1

pranay@localhost:~$ sudo mkfs -t ext4 /dev/share/mylv
```

## 1.14 Format the Logical Volumes

**mkfs -t ext4 /dev/share/mylv**

**mkfs -t ext4 /dev/share/mylv1**

**A:**

```
pranay@localhost:~$ sudo mkfs -t ext4 /dev/share/mylv
mke2fs 1.42.9 (28-Dec-2013)
Filesystem label=
OS type: Linux
Block size=1024 (log=0)
Fragment size=1024 (log=0)
Stride=0 blocks, Stripe width=0 blocks
25688 inodes, 102400 blocks
5120 blocks (5.00%) reserved for the super user
First data block=1
Maximum filesystem blocks=33685504
13 block groups
8192 blocks per group, 8192 fragments per group
1976 inodes per group
Superblock backups stored on blocks:
    8193, 24577, 40961, 57345, 73729

Allocating group tables: done
Writing inode tables: done
Creating journal (4096 blocks): done
Writing superblocks and filesystem accounting information: done

pranay@localhost:~$ sudo mkfs -t ext4 /dev/share/mylv1
mke2fs 1.42.9 (28-Dec-2013)
Filesystem label=
OS type: Linux
Block size=4096 (log=2)
Fragment size=4096 (log=2)
Stride=0 blocks, Stripe width=0 blocks
65536 inodes, 262144 blocks
13107 blocks (5.00%) reserved for the super user
First data block=0
Maximum filesystem blocks=268435456
8 block groups
32768 blocks per group, 32768 fragments per group
```

**1.15 create mount points where you want to mount these LVs**

**1.16 add an appropriate entry to the /etc/fstab file, and mount the filesystem.**

**sudo mkdir -pv/lvmdata**

**sudo mkdir -pv /lvmdata1**

**entries in /etc/fstab file**

**/dev/share/mylv/lvmdata ext4 defaults 1 2**

**/dev/share/mylv1 /lvmdata1 ext4 defaults 1 2**

**run the command "mount -a" to load the mount points**

**1.17 to verify the volume has been created correctly, You can use the commands df, lvs,**

**and vgs**

**1.18 Extend Logical Volume. for ex: For example, to add 500MB more to our LV mylv created from VG share, run the following command:**

**lvextend --size +500M --resizefs share/mylv**

**A: 1.15 to 1.18**

```
[pranay@localhost ~]$ sudo mkdir -pv /lvmdata
mkdir: created directory '/lvmdata'
[pranay@localhost ~]$ sudo mkdir -pv /lvmdata1
mkdir: created directory '/lvmdata1'
[pranay@localhost ~]$ sudo vi /etc/fstab
[pranay@localhost ~]$ mount -a
mount: only root can use "--all" option
[pranay@localhost ~]$ sudo mount -a
[pranay@localhost ~]$ sudo lvs
  LV      VG      Attr      LSize   Pool Origin Data%  Meta%   Move Log Cpy%Sync Convert
  mylv    share -wi-ao---- 100.00m
  mylv1   share -wi-ao----  1.00g
[pranay@localhost ~]$ sudo lvextend --size +500M --resizefs share/mylv
Size of logical volume share/mylv changed from 100.00 MiB (25 extents) to 600.00 MiB (150 extents).
Logical volume share/mylv successfully resized.
resize2fs 1.42.9 (28-Dec-2013)
Filesystem at /dev/mapper/share-mylv is mounted on /lvmdata; on-line resizing required
old_desc_blocks = 1, new_desc_blocks = 5
The filesystem on /dev/mapper/share-mylv is now 614400 blocks long.
```